

# **Report - Variational Sparse Coding**

# Table of contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Datasets . . . . .	3
1.2	Report Structure . . . . .	3
<b>2</b>	<b>Auto Encoder</b>	<b>5</b>
2.1	Architecture . . . . .	5
2.2	Applications . . . . .	5
2.3	Limitations . . . . .	5
<b>3</b>	<b>Variational Auto Encoder</b>	<b>6</b>
3.0.1	VAE's latent space . . . . .	6
3.0.2	Evidence Lower Bound (ELBO) . . . . .	6
3.0.3	$\beta$ -VAE . . . . .	7
<b>4</b>	<b>Variational Sparse Coding</b>	<b>8</b>
4.1	Motivation . . . . .	8
4.2	Recognition Model . . . . .	9
4.3	Training Procedure . . . . .	10
4.3.1	Prior Distribution & Objective . . . . .	10
4.3.2	Warm-Up Strategy . . . . .	11
<b>5</b>	<b>Experiments</b>	<b>12</b>
<b>6</b>	<b>Conclusion</b>	<b>13</b>

# 1 Introduction

Unsupervised learning in high-dimensional data poses significant challenges, particularly in discovering interpretable features and enabling controllable generation.

Variational Autoencoders (VAEs) provide a probabilistic framework for mapping complex data to lower-dimensional latent spaces, but they often fail to disentangle underlying factors of variation, especially when the number of true sources is unknown or when observations exhibit diverse attribute combinations.

The “[Variational Sparse Coding](#)” (VSC) paper by Francesco Tonolini et al. proposes a novel extension of VAEs, incorporating sparsity in the latent space via a Spike and Slab prior to address these issues.

This report for a statistical course explores the VSC model, detailing its theoretical foundations, implementation, and empirical validation.

## 1.1 Datasets

- **MNIST** : TODO...
- **Fashion-MNIST**: A collection of 28x28 grayscale images of clothing items (e.g., T-shirts, trousers), comprising 60,000 training and 10,000 test samples. It serves as a drop-in replacement for MNIST, offering richer variability for testing generative models.
- **Smiley** : TODO...
- **UCI HAR**: A dataset of accelerometer and gyroscope signals from smartphones, capturing six human activities (e.g., walking, sitting) across 30 subjects, with preprocessed segments of 128 time steps.

## 1.2 Report Structure

- **Autoencoders**: Introduces the basics of autoencoders and their role in high-dimensional data analysis.
- **Variational Autoencoders**: Explains VAEs, focusing on the Evidence Lower Bound (ELBO) and latent space regularization.
- **Variational Sparse Coding**: Details the VSC model, including the Spike and Slab prior and warm-up training strategy.

- **Experiments:** Summarizes empirical results, comparing VSC to -VAEs across multiple tasks.
- **Conclusion:** Highlights VSC's contributions and future directions.

## 2 Auto Encoder

An encoder is a neural network that compresses high-dimensional input data into a lower-dimensional latent representation.

### 2.1 Architecture

1. **Encoder:** Maps input  $x \in \mathbb{R}^D$  to latent code  $z \in \mathbb{R}^d$  (compression).
2. **Decoder:** Reconstructs  $\hat{x}$  from  $z$  (reconstruction).

$$\mathcal{L}_{\text{AE}} = \text{Reconstruction\_term} = \|x - \text{Decoder}(\text{Encoder}(x))\|^2$$

Pytorch implementation: `logic/model/autoencoder.py`

### 2.2 Applications

- Feature extraction for clustering/visualization.
- Denoising (e.g., noisy Fashion-MNIST reconstruction).

### 2.3 Limitations

- Deterministic: No probabilistic latent space.
- No control over latent structure (e.g., disentanglement).

## 3 Variational Auto Encoder

Variational Autoencoders (VAEs) extend autoencoders into a probabilistic framework, enabling both representation learning and generative modeling.

Unlike traditional autoencoders, VAEs model the latent variables  $z$  as drawn from a distribution, typically a Gaussian, parameterized by the encoder.

### 3.0.1 VAE's latent space

- The encoder outputs parameters  $\mu$  and  $\sigma^2$ 
  - $q_\phi(z|x) = \mathcal{N}(z; \mu, \sigma^2)$ .
- The decoder generates  $p_\theta(x|z)$ .
- The latent space prior is typically a standard Gaussian,
  - $p(z) = \mathcal{N}(0, I)$ , which constrains the latent space to remain centered around the origin, promoting meaningful structure for visualization and generative tasks.
- The goal is to maximize the marginal likelihood  $p(x)$ , but this is intractable due to the integral:

$$p(x) = \int p_\theta(x|z)p(z)dz$$

### 3.0.2 Evidence Lower Bound (ELBO)

VAEs approximate the marginal likelihood  $p(x)$  using the Evidence Lower Bound (ELBO):

$$\mathcal{L}_{\text{ELBO}} = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{\text{KL}}(q_\phi(z|x) \| p(z))$$

- **Reconstruction Term:**  $\mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)]$ 
  - encourages accurate data reconstruction.
- **KL Divergence Term:**  $D_{\text{KL}}(q_\phi(z|x) \| p(z))$ 
  - regularizes the latent distribution to match the prior.

The ELBO is optimized with respect to encoder parameters  $\phi$  and decoder parameters  $\theta$  using the reparameterization trick for gradient computation.

```
def reparameterize(self, mu: torch.Tensor, logvar: torch.Tensor) -> torch.Tensor:
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)
    return mu + eps * std
```

### 3.0.3 $\beta$ -VAE

The  $\beta$ -VAE introduces a hyperparameter  $\beta$  to control the weight of the KL divergence term in the ELBO:

$$\mathcal{L}_{\beta\text{-VAE}} = \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x|z)] - \beta D_{\text{KL}}(q_{\phi}(z|x) \| p(z))$$

- When  $\beta = 0$ , the  $\beta$ -VAE reduces to a simple autoencoder model.
- When  $\beta = 1$ , the  $\beta$ -VAE reduces to the standard VAE.
- Increasing  $\beta$  encourages disentanglement in the latent space, making it more interpretable but potentially at the cost of reconstruction quality.

The  $\beta$ -VAE is particularly useful in scenarios where interpretability of the latent space is critical, such as in disentangled representation learning. However, careful tuning of  $\beta$  is required to balance reconstruction and disentanglement.

# 4 Variational Sparse Coding

## 4.1 Motivation

What is Posterior Collapse?

- **Problem:** In VAEs, some latent dimensions become “useless” – they encode no meaningful information.
- **Why it happens:**
  - The KL divergence term in ELBO forces latent variables to match the prior.
  - If the decoder is too powerful, it ignores latent variables, leading to **dimensions being permanently inactive**.
  - Result: Model uses only a few dimensions, losing sparsity and disentanglement.

How VSC Fixes It:

### 1. Spike-and-Slab Warm-Up

- **Phase 1** ( $\lambda = 0$ ):
  - Forces latent variables to behave like **binary codes** (spike dominates).
  - Model must “choose” which dimensions to activate (no continuous refinement).
- **Phase 2** ( $\lambda \rightarrow 1$ ):
  - Gradually introduces continuous slab parameters  $(\mu_{i,j}, \sigma_{i,j})$ .
  - Prevents early over-reliance on a few dimensions.

### 2. Sparsity Enforcement

- **KL Sparsity Term:** Penalizes average activation rate  $\bar{\gamma}_u$  if it deviates from  $\alpha$  (e.g.,  $\alpha = 0.01$ ).
- Forces the model to use **only essential dimensions**, avoiding redundancy.

### 3. Dynamic Prior



- Prior  $p_s(z)$  adapts via pseudo-inputs  $x_u$  and classifier  $C_\omega(x_i)$ .
- Prevents trivial alignment with a fixed prior (e.g.,  $\mathcal{N}(0, 1)$ ).

**Result:**

- Latent dimensions stay **sparse and interpretable**.
- No single dimension dominates; features are distributed across active variables. Variational Sparse Coding (VSC) extends VAEs by inducing sparsity in the latent space using a **Spike-and-Slab prior**, enabling feature disentanglement and controlled generation when the number of latent factors is unknown.

## 4.2 Recognition Model

### Spike-and-Slab Encoder Distribution

$$q_\phi(z|x_i) = \prod_{j=1}^J [\gamma_{i,j} \mathcal{N}(z_{i,j}; \mu_{i,j}, \sigma_{i,j}^2) + (1 - \gamma_{i,j}) \delta(z_{i,j})]$$

**Parameters:** All outputs of a neural network (encoder).

- $\gamma_{i,j}$ : Probability that latent dimension  $j$  is *active* for input  $x_i$ .
- $\mu_{i,j}, \sigma_{i,j}$ : Mean and variance of the Gaussian (slab) for active dimensions.
- $\delta(z_{i,j})$ : Dirac delta function (spike) forces inactive dimensions to exactly **0**.

Pytorch implementation of the reparameterization `logic/model/vsc.py`:

```
def reparameterize(self,
    mu: torch.Tensor,
    logvar: torch.Tensor,
    gamma: torch.Tensor
) -> torch.Tensor:

    lamb = self.lambda_val # warm-up factor
    std = torch.exp(0.5 * logvar)
    eps = torch.randn_like(std)

    # Interpolate between a fixed (zero-mean, unit variance) slab
    # and the learned slab.
```

```

slab = lam * mu + eps * (lam * std + (1 - lam))

# Sample binary spike; note: torch.bernoulli is not differentiable.
spike = torch.bernoulli(gamma)

return spike * slab

```

## 4.3 Training Procedure

### 4.3.1 Prior Distribution & Objective

Prior

$$p_s(z) = q_\phi(z|x_{u^*}), \quad u^* = C_\omega(x_i)$$

- **Pseudo-inputs:** Learnable templates  $\{x_u\}$  represent common feature combinations.
- **Classifier:**  $C_\omega(x_i)$  selects the best-matching template  $x_{u^*}$  for input  $x_i$ .

Objective (ELBO with Sparsity)

$$\mathcal{L} = \sum_i \left[ -\text{KL}(q_\phi \| p_s) + \mathbb{E}_{q_\phi} [\log p_\theta(x_i|z)] \right] - J \cdot \text{KL}(\bar{\gamma}_u \| \alpha)$$

- **KL Term:**
  - Aligns encoder  $(\mu_{i,j}, \sigma_{i,j}, \gamma_{i,j})$  with prior  $(\mu_{u^*,j}, \sigma_{u^*,j}, \gamma_{u^*,j})$ .
  - Closed-form formula ensures fast computation.
- **Sparsity Term:**
  - Penalizes deviation from target sparsity  $\alpha$  (e.g., 90% dimensions inactive).

Pytorch implementation in `logic/model/vsc.py`:

```

def compute_sparsity_loss(self, gamma: torch.Tensor) -> torch.Tensor:
    target = torch.full_like(gamma, self.prior_sparsity)
    return nn.functional.binary_cross_entropy(gamma, target)

```

### 4.3.2 Warm-Up Strategy

$$q_{\phi,\lambda}(z|x_i) = \prod_{j=1}^J [\gamma_{i,j} \mathcal{N}(z_{i,j}; \lambda \mu_{i,j}, \lambda \sigma_{i,j}^2 + (1 - \lambda)) + (1 - \gamma_{i,j}) \delta(z_{i,j})]$$

- **Phase 1** ( $\lambda = 0$ ):
  - Slab fixed to  $\mathcal{N}(0, 1)$  (binary-like latent codes).
  - Focus: *Which* features to activate.
- **Phase 2** ( $\lambda \rightarrow 1$ ):
  - Gradually learn  $\mu_{i,j}, \sigma_{i,j}$  (refine *how* to represent features).
- **Avoids collapse**: Prevents premature “freezing” of latent dimensions.

## 5 Experiments

## 6 Conclusion